



Minimum Spanning Trees

- Given an *undirected* graph $G = (V, E)$, with edge costs c_{ij} .
- A spanning tree T of G is a cycle-free subgraph that spans all the nodes.
- The *cost* of T is the *sum* of the costs of the edges in T .
- MST is the smallest cost spanning tree.




Graph G



MST

Applications of MST




- Direct applications: interconnection of entities.
 1. electrical devices (circuit boards)
 2. utilities (gas, oil)
 3. computers or communication devices by high speed lines.
 4. cable service customers
- Indirect applications.
 1. Optimal message passing.
 2. Data storage.
 3. Cluster analysis


Optimality Conditions

- Greedy incremental flavor: add one edge at a time.
- Each step colors an edge of G **blue** (accept) or **red** (reject).
- **Color Invariant (CI):** \exists **MST** containing all blue edges, and no red edges.
- Recall that a cut in $G = (V, E)$ is a partition of its vertices $(X, V - X)$.

Blue and Red Rules



Cut Rule (Blue)




Red Rule (Cycle)

Blue (Cut) Rule: Select a cut not crossed by any blue edge. Among the uncolored edges crossing the cut, make the minimum cost edge blue.


Red (Cycle) Rule: Select a simple cycle with no red edges. Among all uncolored edges of the cycle, make the maximum cost one red.

Generic MST Algorithm

Theorem: *Apply red and blue rules in arbitrary order until neither rule applies. The resulting set of blue edges forms a MST.*



Graph G




MST

- Proof has two parts:
 1. The Color Invariant (CI) holds.
 2. All edges are ultimately colored.


Correctness of Blue Rule

- Let T^* be the MST guaranteed by the CI before the last coloring step.



- Suppose the last step was to color e blue. Consider the cut $(X, V - X)$ to which blue rule applied. Some edge e' of T^* must cross this cut.
- The graph $T^* + e$ contains a cycle containing both e and e' , and $\text{cost}(e) \leq \text{cost}(e')$. (Why?)
- So $T^* + e - e'$ is also a MST.

Correctness of Red Rule



- Suppose edge e colored red. If $e \notin T^*$, then T^* still satisfies CI. Otherwise, consider $T^* - e$. It has two components.
- The cycle used in coloring e has some edge e' with one end in each of these components.
- By choice, $\text{cost}(e') \leq \text{cost}(e)$. Thus, the tree $T^* - e + e'$ is also an MST.

Proof of Completion

- Blue edges form a forest.
- Suppose edge e left uncolored, and neither rule applies.
- If both endpoints of e in same blue tree B , then red rule applies to the cycle $B + e$.
- If endpoints in different blue trees B_1, B_2 , then blue rule applies across the cut separating B_1 and B_2 .
- Thus, the generic MST algorithm is correct.


Prim's Algorithm

- Start at any node s , and set $T = \{s\}$.
- repeat $n - 1$ times


Let T be the current tree.

Choose a minimum cost *uncolored* (u, v) with $u \in T$ and $v \notin T$.

Color (u, v) **blue**, and add v to T .



Graph G





After 3 steps.


Prim's Algorithm

1. **Input:** Graph $G = (V, E)$. For each vertex v , $A(v)$ is list of its neighbors.
2. $key(w)$ is the cost of the cheapest edge (v, w) with $v \in T$. $blue(v, w)$ is the identity of this edge.
3. **Initialize:**
 - $key(v) = \infty$, for all $v \in V$.
 - $H = makeHeap()$.
 - $v = s$.
4. **while** $v \neq NIL$ **do**
 - **for** $w \in A(v)$
 - **if** $cost(v, w) < key(w)$
 - $key(w) = cost(v, w)$
 - $blue(w) = (v, w)$
 - **if** $w \notin H$ $insert(w, H, key(w))$
 - **else** $DecreaseKey(w, H, key(w))$
 - **end for**
 - $v = ExtractMin(H)$.


Illustration



After 3 steps





After 4 steps




Final MST

Illustration



After 3 steps





Blue edges and keys after 3 steps

Kruskal's Algorithm


- Initially make each node of V a singleton tree.
- Scan edges of E in non-decreasing order of cost.
- scan edge e :

If both endpoints of e in the same tree, color it **red**. Otherwise color e **blue**, and merge the two trees.


Illustration



After 4 edges scanned



8, 10 colored red. 12 colored blue




Final MST

Complexity Results



- **Heap implementation of Prim:** $O(m \log_2 n)$.
- **d -Heap:** $O(nd \log_d n + m \log_d n)$.
- **Fibonacci Heap:** $O(m + n \log n)$.
- **Kruskal with Union-Find:**
 $O(m \log n + m \alpha(n))$.
- **Round-Robin:** $O(m \log \log n)$.
- **Latest theoretical bound:** $O(m \alpha(n))$.

Balancing MST and SP



- MST minimizes *total* interconnection cost.
- SPT minimizes *individual* path lengths, from a root source.
- How does each do with the other cost metric?
- In Fig. 2, $d_{mst}(a, c) = 20$, but $d_{spt}(a, c) = 15$.
- In Fig. 3, $cost(spt) = 35$, but $cost(mst) = 30$.

A Pathological Example





- Fig. (i) shows an example where $d_{mst}(a, z) = n$, while $d_{spt} = 2$.
- As $n \rightarrow \infty$, this ratio is unbounded.
- Fig. (ii) shows an example where $cost(spt) = nM$, while $cost(mst) = n - 1 + M$.
- As $M \rightarrow \infty$, this ratio is unbounded.

Balanced Tree Theorem

Theorem: *Pick any constant $\alpha > 1$, and let $\beta = 1 + \frac{2}{\alpha-1}$. Given $G = (V, E)$ and a node s , there always exists a spanning tree T rooted at s such that*

- $d_T(s, v) \leq \alpha \cdot d_{spt}(s, v)$, for any $v \in V$.
- $cost(T) \leq \beta \cdot cost(MST)$.




- Which T for these pathological cases?

Pre-Order Numbering


Pre-Order Lemma: *Let T be a spanning tree with root s . Let z_0, z_1, \dots, z_k be any $k + 1$ nodes listed in their pre-order sequence. Then,*

$$\sum_{i=1}^k d_T(z_{i-1}, z_i) \leq 2\text{cost}(T).$$



- Draw T in the plane. Let W be the doubling “walk” around T (pre-order).
- Each edge visited twice, $\text{cost}(W) = 2\text{cost}(T)$.

Proof




- $d_T(u, v)$ is tree path length.
 - Due to pre-ordering, **first** occurrence of z_{i-1} is **before** that of z_i .
 - Mark off portions of W that join first occurrence of z_{i-1} to that of z_i .
 - No edge of W marked more than once.
- Thus,


$$\sum_{i=1}^k d_T(z_{i-1}, z_i) \leq \text{cost}(W) \leq 2\text{cost}(T).$$

Algorithm

- $\beta = 1 + \frac{2}{\alpha-1}$. E.g. $\alpha = 2, \beta = 3$.
- Compute MST and its pre-order numbers, starting with s .
- Compute $d_{spt}(s, v)$, for all v .
- Initialize $H = MST$.
- for each node v in pre-order do
 if $d_H(s, v) > \alpha \cdot d_{spt}(s, v)$ then
 add to H all edges of path $P_{spt}(s, v)$.
- Output SPT of H , rooted at s .




Graph G




MST


Illustration




MST



(1,5) added



(1,6) added



Final tree T output

Analysis

- T satisfies $d_T(s, v) \leq \alpha \cdot d_{spt}(s, v)$; algorithm adds SP whenever needed.
- Let z_0, z_1, \dots, z_k be the pre-order sequence of vertices that caused SP edges to be added to H .
- When z_i is examined, H contains the shortest path to z_{i-1} . So, we must have

$$d_{spt}(s, z_{i-1}) + d_{mst}(z_{i-1}, z_i) > \alpha \cdot d_{spt}(s, z_i).$$

- Sum these over all i :

$$\begin{aligned} \alpha d_{spt}(s, z_1) - d_{spt}(s, z_0) &< d_{mst}(z_0, z_1) \\ \alpha d_{spt}(s, z_2) - d_{spt}(s, z_1) &< d_{mst}(z_1, z_2) \\ &\vdots \\ \alpha d_{spt}(s, z_k) - d_{spt}(s, z_{k-1}) &< d_{mst}(z_{k-1}, z_k) \end{aligned}$$

- $\sum_{i=1}^k (\alpha - 1) d_{spt}(s, z_i) < \sum_{i=1}^k d_{mst}(z_{i-1}, z_i)$

Analysis

- Thus, we have

$$\begin{aligned}\sum_{i=1}^k (\alpha - 1) d_{spt}(s, z_i) &< \sum_{i=1}^k d_{mst}(z_{i-1}, z_i) \\ &\leq 2 \text{cost}(MST)\end{aligned}$$

- Thus,

$$\sum_{i=1}^k d_{spt}(s, z_i) < \frac{2}{\alpha - 1} \text{cost}(MST)$$

- Since $\text{cost}(H) = \text{cost}(MST) + \sum_{i=1}^k d_{spt}(s, z_i)$, we get

$$\text{cost}(H) < \left(1 + \frac{2}{\alpha - 1}\right) \text{cost}(MST)$$


- Theorem proved.

Graph Spanners


- A subgraph $H = (V, E')$ is a **t -spanner** of graph $G = (V, E)$, where $E' \subseteq E$, if for all pairs $u, v \in V$,

$$\text{dist}_H(u, v) \leq t \cdot \text{dist}_G(u, v)$$

- The parameter t is called the **stretch factor**.
- Total cost of H is the sum of costs of edges in H .
- Determine H with minimum total cost such that stretch factor is at most t .




Graph G



Spanner H

Graph Spanners

- We want a small-size subgraph H that approximates well the shortest paths between all pairs of vertices in G .



- For instance, in building transportation networks, tradeoff between stretch factor and total network cost.
- What's the best H for this example, using Euclidean distances?

Spanner Algorithm

- The following simple, elegant algorithm works.

```
Spanner ( $G, t$ )  
   $H \leftarrow (V, \emptyset);$   
  Sort edges of  $G$  in increasing cost order;  
  Let  $e_1, e_2, \dots, e_m$  be this order;  
  for  $i = 1$  to  $m$  do  
    Let  $e_i = (u, v);$   
    if  $\text{dist}_H(u, v) > t \cdot w(e_i)$   
       $H \leftarrow H \cup \{e_i\};$   
    end  
  end  
  end  
Output  $H;$ 
```

Correctness


Lemma: The subgraph H output by Spanner is a t -spanner of G .

- Let $P(u, v)$ be the shortest path in G .
- Consider an edge e on this path.
- e was either added to H by spanner algorithm, or there is a path of length $t \cdot w(e)$ in H .
- Concatenating these approximate paths in H , we get a path in H whose length is at most $t \cdot P(u, v)$.

Correctness

Lemma: If C is a cycle in H , $\text{size}(C) > t + 1$.

- Suppose \exists a cycle C with $\leq t + 1$ edges.
- Let $e_{\max} = (u, v)$ be its most costly edge.
- When algorithm scans e_{\max} , all other edges of C have been added.




- Thus, there is already a path from u to v of cost $\leq t \cdot e_{\max}$ (each of the t edges is $\leq w(e_{\max})$).
- So, edge e_{\max} will not be added.

Correctness

Lemma: $w(C - e) > t \cdot w(e)$, for any $e \in C$.

- Suppose \exists a cycle C and edge e with $w(C - e) \leq t \cdot w(e)$.



- So, $w(C) \leq (t + 1)w(e) \leq (t + 1)w(e_{\max})$.
- When considering e_{\max} , it won't be added!

Correctness

Lemma: MST contained in H .

- Straightforward. Prove it yourself.
- The key is to show

$$w(E_v) \leq \frac{2w(MST)}{t-1}$$

where E_v is edges of $H - MST$ incident to vertex v .

- This is sufficient to prove

$$\begin{aligned} w(H) &\leq w(MST) + \frac{1}{2} \sum_v w(E_v) \\ &\leq w(MST) \left(1 + \frac{n}{t-1} \right) \end{aligned}$$


Proof

Theorem: $w(E_v) \leq \frac{2w(MST)}{t-1}$.


- Let e_1, e_2, \dots, e_k be edges in E_v .
- For any edge $e \in H$, and **any** path P in H joining the endpoints of e , we have

$$t \cdot w(e) < w(P)$$

- Let **Poly** be the polygon defined by doubling MST.



Proof



- W_i is perimeter of Poly after e_1, \dots, e_i added. Initially, $W_0 = MST$.
- What happens when $e_i = (u, v)$ is added?

$$\begin{aligned} W_i &= W_{i-1} + w(e_i) - w(P(u, v)) \\ &< W_{i-1} - w(e_i) \cdot (t - 1) \end{aligned}$$

$$(t - 1)w(e_i) < W_{i-1} - W_i$$

- Thus, $(t - 1) \sum_{i=1}^k w(e_i) \leq W_0 - W_k$.
- Since $W_0 = 2w(MST)$ and $W_k > 0$,

$$E_v = \sum_{i=1}^k w(e_i) < \frac{2w(MST)}{t - 1}.$$

Spanner Results

- Althofer et al. result gives spanners with $O(n^{1+\frac{2}{t-1}})$ edges.
- Peleg-Schaffer shows that there are n -vertex graphs, whose all t -spanners contain $\Omega(n^{1+\frac{1}{t}})$ edges.
- Much better results know if distances are Euclidean.
- In Euclidean k -space (constant k) spanners with $O(n)$ edges exist for any constant stretch factors t .
- Such spanners possible with $O(MST)$ weight, $O(1)$ degree, $O(\log n)$ diameter.