

Stack Overflow is a community of 4.7 million programmers, just like you, helping each other. Join them, it only takes a minute:

Sign up

x

How to choose between map and unordered_map?

Suppose I wanted to map data with a string as the key. What container should I have chosen, `map` or `unordered_map` ? `unordered_map` takes up more memory so let's suppose memory isn't an issue, and the concern is speed.

`unordered_map` should generally give average complexity of $O(1)$ with the worst case of $O(n)$. In what cases would it get to $O(n)$? When does a `map` get more time efficient than `unordered_map` ? Does it happen when n is small?

Assuming I would use STL `unordered_map` with the default haser Vs. `map`. string is the key.

If I'm going to iterate over the elements rather than access an individual element each time, should I prefer `map` ?

c++

dictionary

data-structures

stl

unordered-map

edited Jul 31 at 5:07

 Krishna_Oza

396218

asked Dec 10 '12 at 10:58

 StackHeapCollision

4451315

- 2 Do you need to items in the mapping to be sorted? – Joachim Pileborg Dec 10 '12 at 11:02
- Which implementation of `unordered_map` uses more memory? – Peter Wood Dec 10 '12 at 11:07
- You always have memory overhead in a hash map, although it is typically negligible. – ypnos Dec 10 '12 at 11:09
- Yes items would be stored and assuming I would use STL `unordered_map` with the default haser Vs. `map`. – StackHeapCollision Dec 10 '12 at 11:21
- It's a minor point but as you mention iteration, it's worth pointing out that if you iterate while inserting elements, you should favor `map` over `unordered_map`. – JMcF Aug 29 '13 at 23:51


4 Answers

In practice, if memory is no issue, `unordered_map` is always faster if you want single element access.

The worst case is theoretical and bound to a single hash accounting for all of the elements. This is not of practical relevance. The `unordered_map` gets slower as soon as you have at least $\log N$ elements belonging to the same hash. This is also not of practical relevance. In some special scenarios you could use a specific hashing algorithm that ensures a more uniform distribution. For ordinary strings that don't share a specific pattern, the generic hash functions coming with `unordered_map` are just as good.

If you want to traverse the map (using iterators) in a sorted fashion, you cannot use `unordered_map`.

answered Dec 10 '12 at 11:07

 ypnos

25.9k65694

	map	unordered_map
element ordering	strict weak	n/a
common implementation	balanced tree or red-black tree	hash table
search time	$\log(n)$	$O(1)$ if there are no hash collisions Up to $O(n)$ if there are hash collisions $O(n)$ when hash is the same for any key
Insertion time	$\log(n)$ +rebalance	Same as search
Deletion time	$\log(n)$ +rebalance	Same as search

needs comparators	only <code>operator <</code>	only <code>operator ==</code>
needs hash <code>function</code>	<code>no</code>	yes
common <code>use case</code>	<code>when</code> good hash <code>is</code> <code>not</code> possible <code>or</code> too slow. <code>Or when</code> order <code>is</code> required	In most other cases.

edited Dec 10 '12 at 11:26

answered Dec 10 '12 at 11:16

user1773602

This is exactly what I want! Thanks! – [daizuozhuo](#) Nov 12 '13 at 6:42

Comment about common implementation: A red-black tree is a *kind* of balanced tree (or more specifically, a kind of self-balancing binary search tree). – [HelloGoodbye](#) Oct 18 at 14:44

In what cases would it get to $O(n)$?

if you have such a **bad** hash function which produces the same hash value for all input strings (i.e. produce collisions)...

What container should I have chosen, map or unordered_map?

It is always the questions of requirements and kind/amount of data do you have.

When does a map get more time efficient than unordered_map?

It is just different structures. You'd better to make a choice to use one of them depending on your typical use cases (taking in account what kind of data do you have and its amount)

Does it happen when n is small?

In case of small data amount everything depends on particular STL implementation... So sometimes even a plain vector/array could be faster than associative containers...

answered Dec 10 '12 at 11:07



zaufi

2,767

6

17

What container should I have chosen, map or unordered_map? unordered_map takes up more memory so let's suppose memory isn't an issue, and the concern is speed.

Profile and then decide. unordered_map is generally faster, but it varies per case.

In what cases would it get to $O(n)$?

When the hashing isn't good and a bunch of elements are being assigned to the same bins.

When does a map get more time efficient than unordered_map? Does it happen when n is small?

Probably not, but profile it if you really care. Having a container with a small size be the bottleneck of your program seems extremely unlikely. Anyway, a simple vector with linear search may be faster for such cases.

The most important thing when deciding is the requirements of ordering and lack of iterator invalidation. If you need either, you pretty much have to use map. Otherwise, unordered_map.

answered Dec 10 '12 at 11:09



Pubby

31.7k

3

70

130

