

## String to Number Conversion in C Takes its Toll

Converting a string to a number in C is no simple affair. Many of you may have heard of `atoi(3)`; one of the ways to convert a string to a number. Unfortunately, modern thinking says that `atoi` should never be used, and so its use is **discouraged**. Instead, we now have `strtol`. <sup>1</sup>

### What's wrong with `atoi`?

The reason for discouraging use of `atoi` stems from the fact that there is no way to detect if overflow or underflow has occurred, and no way to check if the entire string has been converted (aka there's no way to detect `atoi("123garbage")`). Consider the following code:

```
// 2^32+1, assuming 32-bit int
const char *uintmax_plus_one = "4294967297";
printf("%d\n", atoi(uintmax_plus_one));
```

When run instead of printing 4294967296 as expected, this program will print 1! The vast majority of programs do not check for or properly handle this case, and so you can end up with situations like the following:

```
// 2^32+1, assuming 32-bit int
const char *uintmax_plus_one = "4294967297";
malloc(atoi(uintmax_plus_one));
```

Now we're allocating far less memory than we expected. This problem can quickly become the source of an **integer overflow** vulnerability.

### Enter `strtol`!

In order to do the conversion safely, we instead should use `strtol`. It is unfortunately quite difficult to call this function properly. Consider the following documentation pulled from the BSD Library Functions Manual's section on `strtol`:

The `strtol()`, `strtoll()`, `strtoimax()`, and `strtobq()` functions return the result of the conversion, unless the value would underflow or overflow. If no conversion could be

performed, 0 is returned and the global variable `errno` is set to `EINVAL` (the last feature is not portable across all platforms). If an overflow or underflow occurs, `errno` is set to `ERANGE` and the function return value is clamped according to the following table.

Function	underflow	overflow
<code>strtol()</code>	<code>LONG_MIN</code>	<code>LONG_MAX</code>
<code>strtoll()</code>	<code>LLONG_MIN</code>	<code>LLONG_MAX</code>
<code>strtoimax()</code>	<code>INTMAX_MIN</code>	<code>INTMAX_MAX</code>
<code>strtoq()</code>	<code>LLONG_MIN</code>	<code>LLONG_MAX</code>

Based on this, the two ways to check for an overflow are to check if `strtol` returns 0 or to check if `errno` is set to `ERANGE`. There's another simple case where `strtol` returns 0 specifically if the input string to `strtol` is `"0"`, so in order to accurately detect range errors, we must check for `ERANGE`. This isn't quite so simple either however, as if no error has occurred, `strtol` will not change the value of `errno`. If this happens, and some previous code has set `errno` to `ERANGE` you will erroneously think that a range error has occurred. So now in order to check for range errors you must reset `errno` to a value that indicates that no error has occurred. Now, what value is that? Thankfully, [POSIX.1-2008](#) has considered this possibility, and defined that "No function in this volume of POSIX.1-2008 shall set `errno` to 0," meaning that no error in all of POSIX will have the value 0. So now we can do the following:

```

errno = 0;
long val = strtol(str, NULL, 10);
if (errno == ERANGE) {
    switch(val) {
        case LONG_MIN:
            // underflow
            break;
        case LONG_MAX:
            // overflow
            break;
        default:
            assert(false); // impossible
    }
} else if (errno != 0) {
    // something else happened. die die die
}

```

But wait, there's more! This works fine for detecting range errors, but fails to detect garbage at the end of the string. Thankfully, `strtol` lets us handle this too via its `char **endptr` argument.

If `endptr` is not `NULL`, `strtol()` stores the address of the first invalid character in `*endptr`. If there were no digits at all, however, `strtol()` stores the original value of `str` in `*endptr`. (Thus, if `*str` is not `'\0'` but `**endptr` is `'\0'` on return, the entire string was valid.)

Using this, we can detect if there was garbage at the end of the string by passing in a `char **` value for `endptr`.

```
errno = 0;
char *endptr;
long val = strtol(str, &endptr, 10);
if (errno == ERANGE) {
    switch(val) {
        case LONG_MIN:
            // underflow
            break;
        case LONG_MAX:
            // overflow
            break;
        default:
            assert(false); // impossible
    }
} else if (errno != 0) {
    // something else happened. die die die
} else if (*endptr != '\0') {
    // garbage at end of string
}
```

And now we've turned a relatively simple one-line call to `atoi` into twenty lines of code.


**FML**

## Enter `strtonum`

The great folks over at **OpenBSD** have made a nice replacement for `strtol` which fixes all of the issues discussed. Called **`strtonum`**, the function doesn't allow trailing characters and makes it easy to determine if a range error has occurred. The one drawback is that `strtonum` is an OpenBSD extension, and so is not found in any standard. If you want to use `strtonum` on other platforms, you can grab the source [here](#).

1.

`strtol`, short for "string to long" is only one of a few of such functions for converting from a string to a number. Also in this family are `strtoimax`, `strtoll` and `strtouq`. ↩

 Recommend Share

Sort by Best ▾

[Join the discussion...](#)**Droggl** • 3 years ago

I wondered for a long time why the heck we need exception handling at all and whats wrong with using the good ole C error handling style. This post illustrates it pretty well ;)

On a more serious note: I always wondered why atoi() was deprecated, thanks for the thorough explanation!

3 ^ | v • Reply • Share ›

**jean luc** • a year ago

I wish something like this was posted on reference sites for all functions. It would help more people write safer code.

^ | v • Reply • Share ›

**William Orr** • a year ago

Any thoughts on NetBSD's strtol(3)?

^ | v • Reply • Share ›

**Jayson** • 3 years ago

Not that it changes the intent of the post, but according to the text you posted from the manual "If no conversion could be performed, 0 is returned AND the global variable errno is set to EINVAL". You should have a check for errno == EINVAL.

^ | v • Reply • Share ›

**eatnumber1** Mod → Jayson • 3 years ago

That's covered by the errno != 0 portion. I didn't call it out explicitly like I did with ERANGE because I didn't explicitly discuss this case like I did with range errors and garbage at the end of the string.

^ | v • Reply • Share ›

**ikk** • 3 years ago

It seems like every line in C turns into 20 lines if you want to do it correctly. Convert a number to string? Check. Allocate or reallocate a block of memory? Check. Make a fool-proof call to snprintf, str(n)cat, strftime, scanf or whatever? Check.

If C functions could do the right thing by default I'd be so happy!




^ | v • Reply • Share ›

**John** → ikk • 3 years ago

<http://www.jwz.org/doc/wors...>

^ | v • Reply • Share ›

---

 [Subscribe](#)  [Add Disqus to your site](#) [Add Disqus](#) [Add](#)  [Privacy](#)

*Content by* [RUSSELL HARMON](#). *Design by* [MARK REID](#)  
([SOME RIGHTS RESERVED](#))

*Powered by* [JEKYLL](#). *Hosted by* [GITHUB](#).