


# B+ trees

## Basics 3

B+ trees as indexes into database tables

Douglas H. Fisher



1) At **depth 2**, a B+ tree of **order 2** has a **MAXIMUM** of  $5*5*4 = 100$  **items** across all leaves (and a minimum of  $2*3*2 = 12$  items at leaves);

2) At **depth M**, a B+ tree of **order 2** has a **MAXIMUM** of  $5^M * 4$  **items** across all leaves

...

3) At depth M, a B+ tree of order 50 has a **MAXIMUM** of  $(2*50+1)^M * (2*50) = 101^M * 100 > 100^{M+1}$  **items** at each leaf.

## Two questions

What is the **MAXIMUM number of items** that a B+ tree of **order 50 and depth 2** can have at its leaves?

What is the **MINIMUM number of links** that would have to be followed to find an item at the leaf of a **B+ tree of order 50 and that contained 1,006,201 items** across all leaves?

What is the **MAXIMUM** number of items that a B+ tree of **order 50 and depth 2** can have at its leaves? **1,020,100**

The maximum number of child links of each internal node is  $2*50+1$  or 101. The maximum number of items at each leaf is 100.

- 101 child links (max) at the root (depth 0)
- $101^2$  (= 10,201) child links (max) across all depth 1 nodes
- $101^2 * 100$  (= 1,020,100) items (max) across all depth 2 leaves

What is the **MINIMUM** number of pointers that would have to be followed to find an item at the leaf of a **B+ tree of order 50 and that contained 1,006,201 items** across all leaves? **2 pointers at minimum**


A depth 1 B+ tree of order 50 could hold at most  $101*100 = 10,100$  (not enough), but a depth 2 tree has the capacity (see above)

## B+ trees and binary search trees

- 1) The depth of a B+ tree of order 50 that holds  $N$  items at leaves is  $O(\log_{100} N)$ . The depth also corresponds to the number of pointers that would have to be followed to find an item.
- 2) In contrast, the depth of a good ol' suitably balanced binary search tree is  $O(\log_2 N)$ . Again, depth corresponds to the number of pointers that would have to be followed.
- 3) If  $N$  is 1,000,000, then 2-3 pointers must be followed in B+ tree of order 50 (with no less than 50 and up to 100 keys in each node)
- 4) For a balanced binary search tree, there would be about 20 pointers followed in  $N$  is 1,000,000
- 5) Looks like a big win for the B+ tree, **but not so fast!**


... but not so fast!

For EACH node of the B+ tree, we must find where the search is to be directed. We could use sequential search for B+ trees of small orders, but for anything but the smallest orders, **binary search** would be better (and the average cost of binary search is almost the worst case of about  $\log_2 n$ )

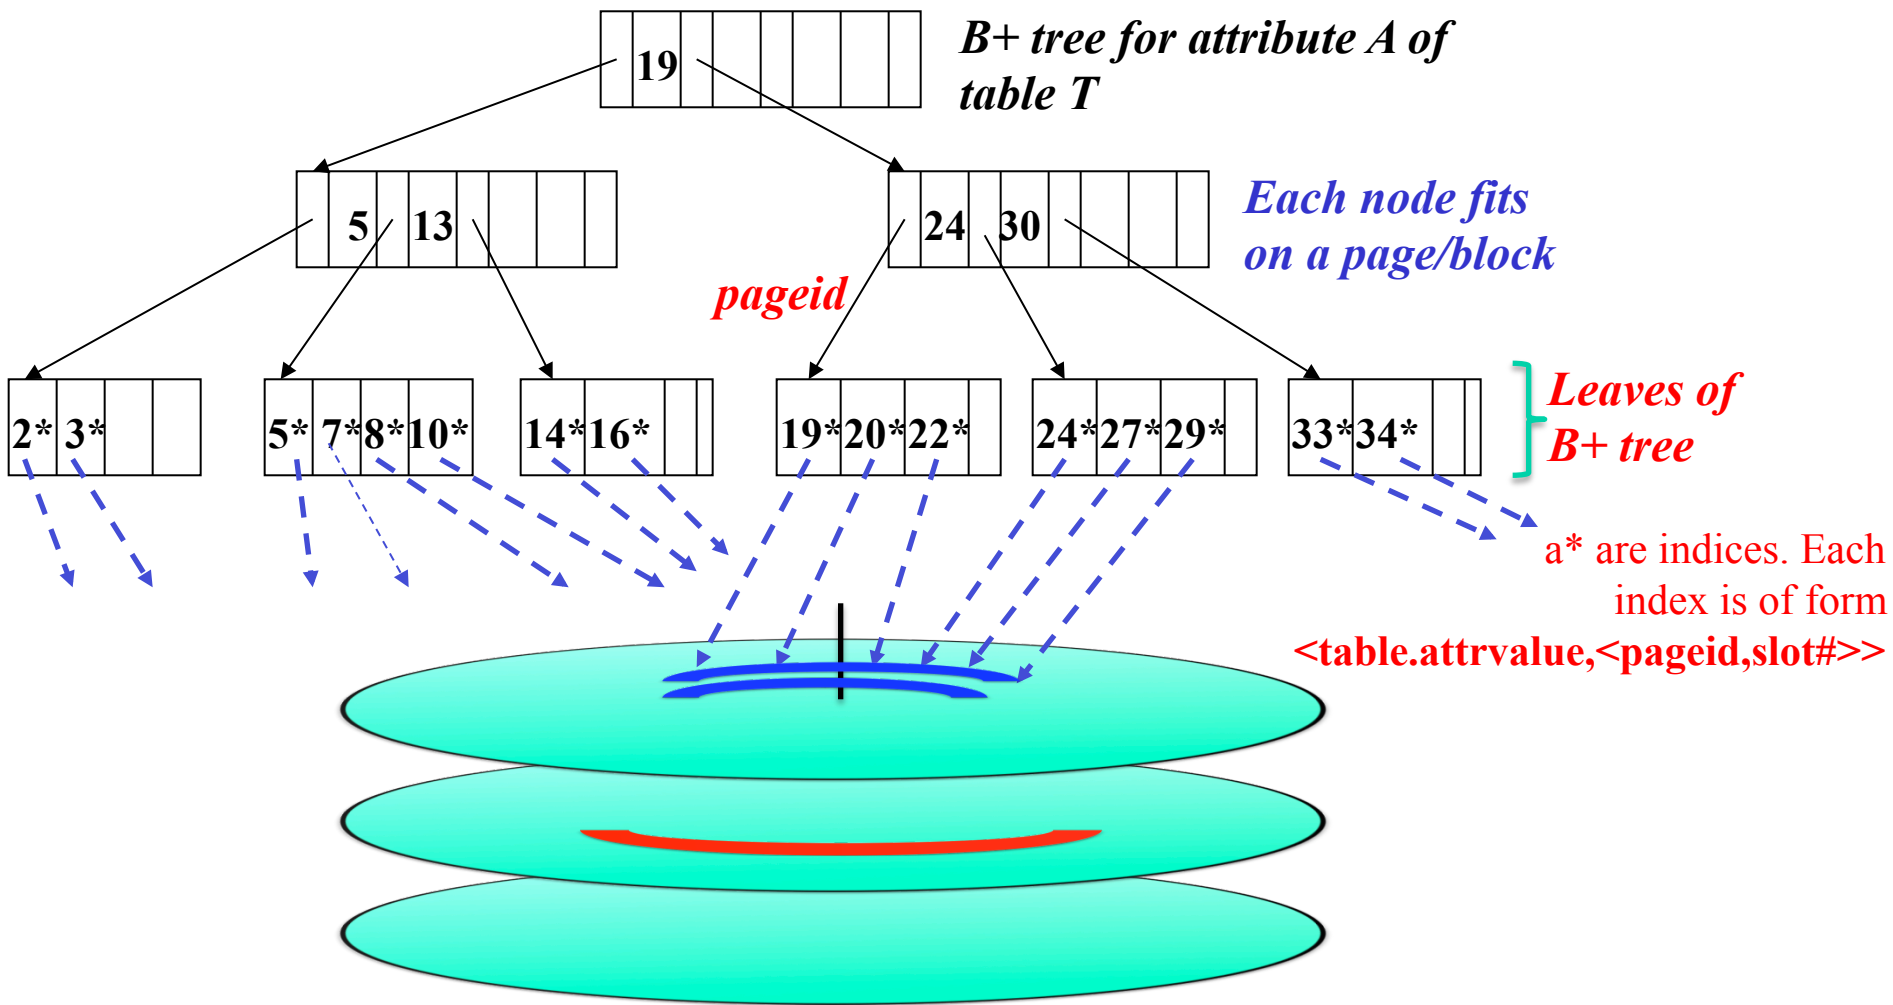


So the # “steps” using a B+ tree of order 50 (100 entries)  
= number of nodes times cost per node  
 $\approx \log_{100} n * \log_2 100$  (generalizes to any order)  
=  $\log_2 n$   
 $\approx$  # “steps” using a balanced binary search tree

- 1) The asymptotic cost of using a B+ tree and a binary search tree are the same,**
- 2) but in accessing database tables, following pointers is VERY expensive,**
- 3) because database tables (and the B+ trees that index them) are stored on disk,**
- 4) where each access takes time proportional to  $\text{disk rotational delay} * \text{seek time} * \text{read(write) time}$**
- 5) so we want to minimize the # of pointers followed**



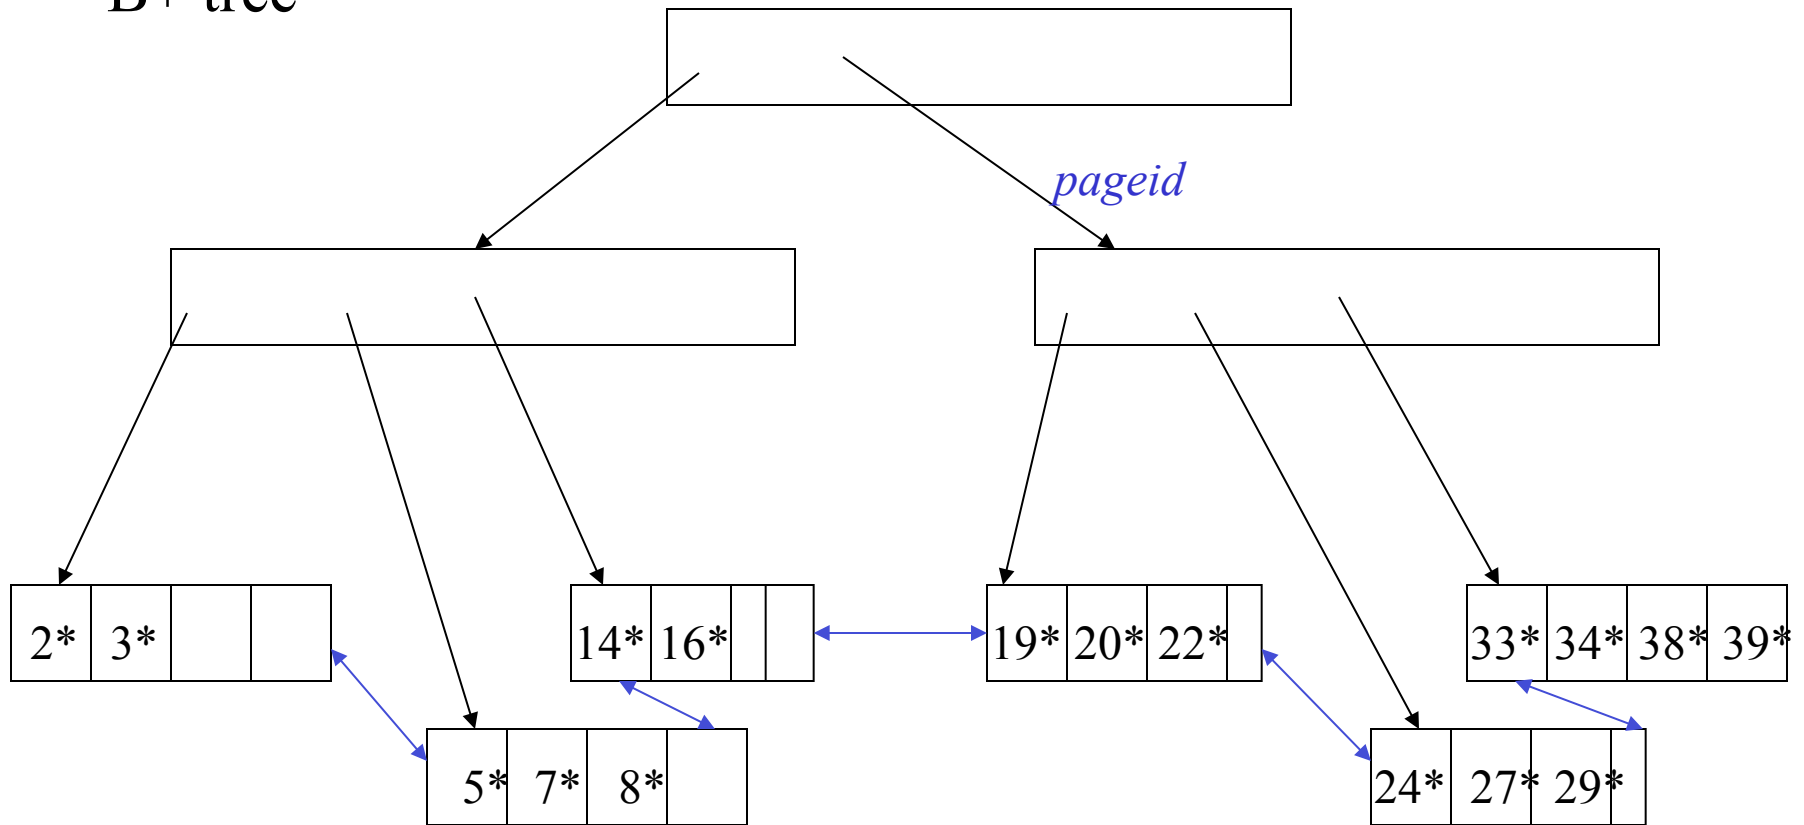




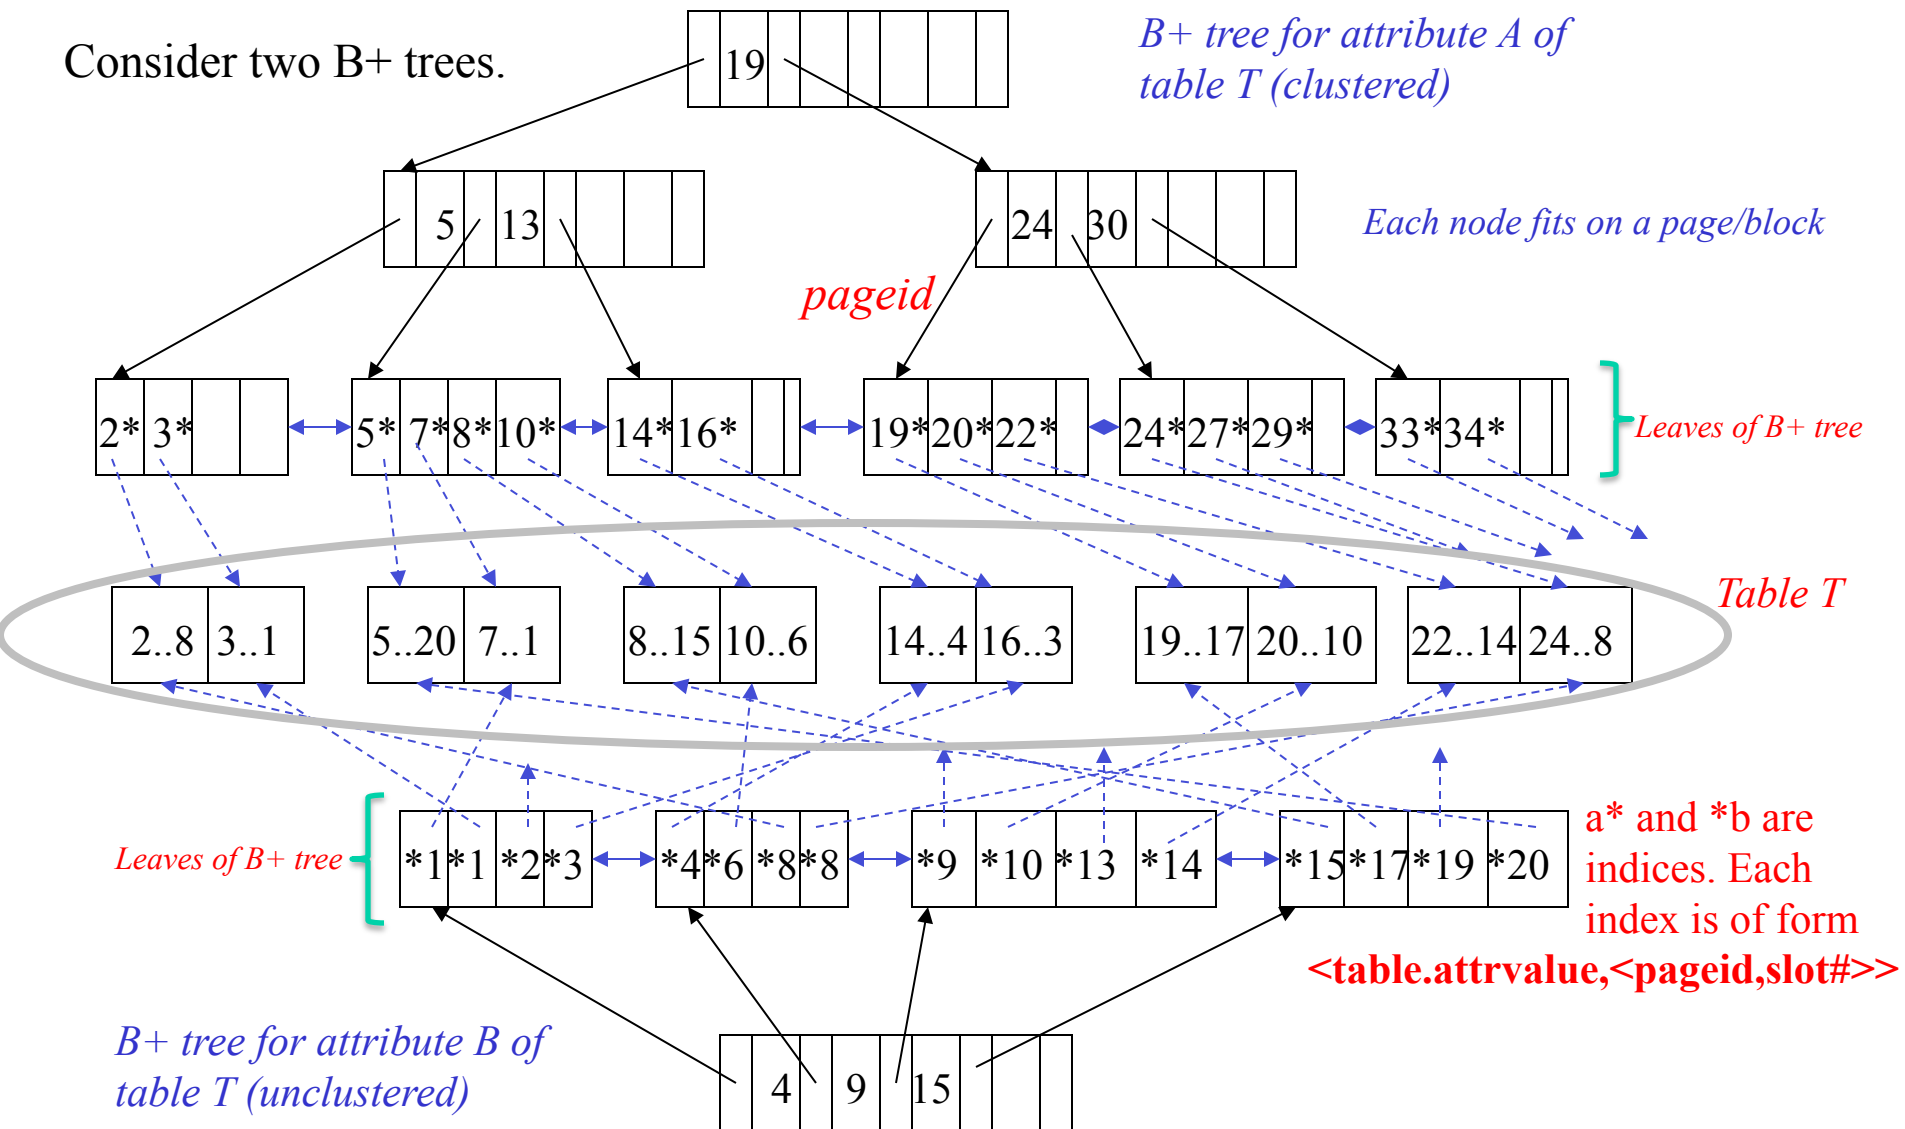
As an example, suppose that a block (or page) holds  $2^{12}$  bytes  
 Suppose each tuple of table T requires  $2^4$  bytes  
 Suppose each index for table T requires  $2^3$  bytes  
 B+ tree order  $2^8$  would hold up to  $2^9$  indexes

# B+ tree

*Each node fits on a page/block*



Consider two B+ trees.



**SELECT T.C FROM T WHERE T.A > 14 AND T.B <= 10**

**Exploiting T.A clustered B+ tree index will result in fewer pages being read from disk.**